

LSTNet-For-Cryptocurrency-Market-Prediction

June 11, 2020

0.1 1. Setup

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.graph_objs as go
from plotly.offline import download_plotlyjs, init_notebook_mode, plot, iplot
from tqdm.notebook import tqdm
from datetime import datetime, timedelta

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader

[3]: device = "cuda:0" if torch.cuda.is_available() else "cpu"
device = torch.device(device)
```

0.2 2. Load data

```
[4]: data_path = "Data/BNB_BTC_1h.csv"
data = pd.read_csv(data_path)
print(data.shape)
data.head()
```

(21461, 12)

```
[4]:
```

	Open-Time	Open	High	Low	Close	Volume	\
0	2018-01-01 00:00:00	0.000623	0.000625	0.000614	0.000618	130791.0	
1	2018-01-01 01:00:00	0.000618	0.000625	0.000611	0.000622	89072.0	
2	2018-01-01 02:00:00	0.000622	0.000625	0.000620	0.000622	101296.0	
3	2018-01-01 03:00:00	0.000622	0.000624	0.000619	0.000622	83701.0	
4	2018-01-01 04:00:00	0.000622	0.000628	0.000622	0.000625	103000.0	

	Close-Time	Quote-Asset-Value	Number-of-Trades	\
0	2018-01-01 00:59:59	80.955613	3328.0	

1	2018-01-01 01:59:59	55.013318	2653.0
2	2018-01-01 02:59:59	63.078538	2663.0
3	2018-01-01 03:59:59	52.015837	2726.0
4	2018-01-01 04:59:59	64.458631	2921.0

	Taker-Buy-Base-Asset-Volume	Taker-Buy-Quote-Asset-Volume	Ignore
0	54856.0	33.995884	0.0
1	39926.0	24.693081	0.0
2	59607.0	37.145069	0.0
3	53257.0	33.115652	0.0
4	55663.0	34.841113	0.0

```
[5]: columns = data.columns[:7]
data = data[columns]
data.head()
```

```
[5]:
```

	Open-Time	Open	High	Low	Close	Volume \
0	2018-01-01 00:00:00	0.000623	0.000625	0.000614	0.000618	130791.0
1	2018-01-01 01:00:00	0.000618	0.000625	0.000611	0.000622	89072.0
2	2018-01-01 02:00:00	0.000622	0.000625	0.000620	0.000622	101296.0
3	2018-01-01 03:00:00	0.000622	0.000624	0.000619	0.000622	83701.0
4	2018-01-01 04:00:00	0.000622	0.000628	0.000622	0.000625	103000.0

	Close-Time
0	2018-01-01 00:59:59
1	2018-01-01 01:59:59
2	2018-01-01 02:59:59
3	2018-01-01 03:59:59
4	2018-01-01 04:59:59

```
[6]: data.describe()
```

```
[6]:
```

	Open	High	Low	Close	Volume
count	21320.000000	21320.000000	21320.000000	21320.000000	2.132000e+04
mean	0.002113	0.002125	0.002100	0.002113	7.988280e+04
std	0.000804	0.000809	0.000799	0.000804	9.706676e+04
min	0.000567	0.000570	0.000529	0.000568	0.000000e+00
25%	0.001540	0.001549	0.001530	0.001540	3.322559e+04
50%	0.001935	0.001944	0.001926	0.001935	5.791882e+04
75%	0.002371	0.002385	0.002356	0.002371	9.633458e+04
max	0.004738	0.004813	0.004680	0.004736	5.178274e+06

0.3 3. Clean data

0.3.1 3.1 Deal NaN

```
[7]: na_mask = data.isna().any(axis=1)
      print(na_mask.sum())
      data[na_mask].head()
```

141

```
[7]:
```

		Open-Time	Open	High	Low	Close	Volume	Close-Time
76	2018-01-04	04:00:00	NaN	NaN	NaN	NaN	NaN	2018-01-04 04:59:59
913	2018-02-08	01:00:00	NaN	NaN	NaN	NaN	NaN	2018-02-08 01:59:59
914	2018-02-08	02:00:00	NaN	NaN	NaN	NaN	NaN	2018-02-08 02:59:59
915	2018-02-08	03:00:00	NaN	NaN	NaN	NaN	NaN	2018-02-08 03:59:59
916	2018-02-08	04:00:00	NaN	NaN	NaN	NaN	NaN	2018-02-08 04:59:59

As 141 is very small compared to 21461 and DL being robust to such noise, I am just filling Nan with 'ffill'

```
[8]: data = data.ffill()
      print(data.isna().any(axis=1).sum())
      data[na_mask].head()
```

0

```
[8]:
```

		Open-Time	Open	High	Low	Close	Volume	\
76	2018-01-04	04:00:00	0.000612	0.000612	0.000611	0.000612	130.0	
913	2018-02-08	01:00:00	0.001064	0.001073	0.001059	0.001063	42985.9	
914	2018-02-08	02:00:00	0.001064	0.001073	0.001059	0.001063	42985.9	
915	2018-02-08	03:00:00	0.001064	0.001073	0.001059	0.001063	42985.9	
916	2018-02-08	04:00:00	0.001064	0.001073	0.001059	0.001063	42985.9	

		Close-Time
76	2018-01-04	04:59:59
913	2018-02-08	01:59:59
914	2018-02-08	02:59:59
915	2018-02-08	03:59:59
916	2018-02-08	04:59:59

0.3.2 3.2 Clean sequence

```
[9]: data['Open-Time'] = pd.to_datetime(data['Open-Time'])
      data = data.sort_values(by=['Open-Time'])
      start_dt = data.loc[0, 'Open-Time']
      end_dt = data.loc[len(data)-1, 'Open-Time']
      expected_end_dt = start_dt + timedelta(hours=len(data) - 1)
      print(start_dt)
      print(expected_end_dt)
```

```
print(end_dt)
```

```
2018-01-01 00:00:00
2020-06-13 04:00:00
2020-06-11 09:00:00
```

```
[10]: print(len(data))
      print(data['Open-Time'].nunique())
```

```
21461
21461
```

There are no duplicates. But end date expectation is not matching. Lets dig deeper

```
[11]: data['Open-Time-Shifted'] = data['Open-Time'].shift(periods=1)
      data.loc[:, 'Seconds-Lag'] = (data['Open-Time'] - data['Open-Time-Shifted'])[1:
      ↪].apply(timedelta.total_seconds)
      data.head()
```

```
[11]:
```

	Open-Time	Open	High	Low	Close	Volume	\
0	2018-01-01 00:00:00	0.000623	0.000625	0.000614	0.000618	130791.0	
1	2018-01-01 01:00:00	0.000618	0.000625	0.000611	0.000622	89072.0	
2	2018-01-01 02:00:00	0.000622	0.000625	0.000620	0.000622	101296.0	
3	2018-01-01 03:00:00	0.000622	0.000624	0.000619	0.000622	83701.0	
4	2018-01-01 04:00:00	0.000622	0.000628	0.000622	0.000625	103000.0	

	Close-Time	Open-Time-Shifted	Seconds-Lag
0	2018-01-01 00:59:59	NaT	NaN
1	2018-01-01 01:59:59	2018-01-01 00:00:00	3600.0
2	2018-01-01 02:59:59	2018-01-01 01:00:00	3600.0
3	2018-01-01 03:59:59	2018-01-01 02:00:00	3600.0
4	2018-01-01 04:59:59	2018-01-01 03:00:00	3600.0

```
[12]: data['Seconds-Lag'].describe()
```

```
[12]:
```

count	21460.000000
mean	3592.786580
std	113.919953
min	1694.000000
25%	3600.000000
50%	3600.000000
75%	3600.000000
max	3600.000000

Name: Seconds-Lag, dtype: float64

```
[13]: mask = data['Seconds-Lag'] < 3600
      print(mask.sum())
      data[mask].head(n=10)
```

```
[13]:
```

		Open-Time	Open	High	Low	Close	Volume \
946	2018-02-09	09:28:14	0.001063	0.001110	0.001000	0.001080	264762.75
947	2018-02-09	10:00:00	0.001063	0.001110	0.001000	0.001080	264762.75
948	2018-02-09	10:28:14	0.001080	0.001080	0.001050	0.001066	247583.54
949	2018-02-09	11:00:00	0.001080	0.001080	0.001050	0.001066	247583.54
950	2018-02-09	11:28:14	0.001066	0.001078	0.001050	0.001077	249736.63
951	2018-02-09	12:00:00	0.001066	0.001078	0.001050	0.001077	249736.63
952	2018-02-09	12:28:14	0.001076	0.001082	0.001072	0.001078	266598.78
953	2018-02-09	13:00:00	0.001076	0.001082	0.001072	0.001078	266598.78
954	2018-02-09	13:28:14	0.001079	0.001136	0.001078	0.001107	475272.41
955	2018-02-09	14:00:00	0.001079	0.001136	0.001078	0.001107	475272.41

		Close-Time	Open-Time-Shifted	Seconds-Lag
946	2018-02-09	10:28:14	2018-02-09 09:00:00	1694.0
947	2018-02-09	10:59:59	2018-02-09 09:28:14	1906.0
948	2018-02-09	11:28:14	2018-02-09 10:00:00	1694.0
949	2018-02-09	11:59:59	2018-02-09 10:28:14	1906.0
950	2018-02-09	12:28:14	2018-02-09 11:00:00	1694.0
951	2018-02-09	12:59:59	2018-02-09 11:28:14	1906.0
952	2018-02-09	13:28:14	2018-02-09 12:00:00	1694.0
953	2018-02-09	13:59:59	2018-02-09 12:28:14	1906.0
954	2018-02-09	14:28:14	2018-02-09 13:00:00	1694.0
955	2018-02-09	14:59:59	2018-02-09 13:28:14	1906.0

```
[14]: irregular_time_mask = data['Open-Time'].dt.minute != 0
print(irregular_time_mask.sum())
print(data.shape)
data = data[np.logical_not(irregular_time_mask)]
data = data.reset_index(drop=True)
print(data.shape)
```

43

(21461, 9)

(21418, 9)

```
[15]: data['Open-Time'] = pd.to_datetime(data['Open-Time'])
data = data.sort_values(by=['Open-Time'])
start_dt = data.loc[0, 'Open-Time']
end_dt = data.loc[len(data)-1, 'Open-Time']
expected_end_dt = start_dt + timedelta(hours=len(data) - 1)
print(start_dt)
print(expected_end_dt)
print(end_dt)
```

2018-01-01 00:00:00

2020-06-11 09:00:00

2020-06-11 09:00:00

```
[16]: columns = data.columns[:6]
data = data[columns]
print(data.shape)
data.head()
```

(21418, 6)

```
[16]:
```

	Open-Time	Open	High	Low	Close	Volume
0	2018-01-01 00:00:00	0.000623	0.000625	0.000614	0.000618	130791.0
1	2018-01-01 01:00:00	0.000618	0.000625	0.000611	0.000622	89072.0
2	2018-01-01 02:00:00	0.000622	0.000625	0.000620	0.000622	101296.0
3	2018-01-01 03:00:00	0.000622	0.000624	0.000619	0.000622	83701.0
4	2018-01-01 04:00:00	0.000622	0.000628	0.000622	0.000625	103000.0

0.4 4. Define train dataset

```
[17]: train_max_open_time = datetime.strptime('2020-05-31 23:00:00', '%Y-%m-%d %H:%M:
      ↪%S')
train_data = data[data['Open-Time'] <= train_max_open_time].copy()
print(len(data), len(train_data))
train_data.tail()
```

21418 21168

```
[17]:
```

	Open-Time	Open	High	Low	Close	Volume
21163	2020-05-31 19:00:00	0.001827	0.001830	0.001825	0.001830	15025.47
21164	2020-05-31 20:00:00	0.001830	0.001830	0.001822	0.001829	21273.65
21165	2020-05-31 21:00:00	0.001829	0.001829	0.001801	0.001817	48774.92
21166	2020-05-31 22:00:00	0.001817	0.001826	0.001811	0.001812	42385.93
21167	2020-05-31 23:00:00	0.001813	0.001816	0.001793	0.001808	81678.70

```
[18]: min_price = train_data['Low'].min()
max_price = train_data['High'].max()
min_volume = train_data['Volume'].min()
max_volume = train_data['Volume'].max()
print(min_price, max_price)
print(min_volume, max_volume)
```

0.000529 0.0048133999999999999
0.0 5178273.91

```
[19]: train_data.loc[:, 'Open-N'] = (train_data['Open'] - min_price) / (max_price-
      ↪min_price)
train_data.loc[:, 'High-N'] = (train_data['High'] - min_price) / (max_price-
      ↪min_price)
```

```

train_data.loc[:, 'Low-N'] = (train_data['Low'] - min_price) / (max_price-
↪min_price)
train_data.loc[:, 'Close-N'] = (train_data['Close'] - min_price) / (max_price-
↪min_price)
train_data.loc[:, 'Volume-N'] = (train_data['Volume'] - min_volume) /
↪(max_volume - min_volume)

```

```

[20]: class MarketDataset(Dataset):

    def __init__(self, data, history_len):
        self.data = data
        self.history_len = history_len

    def __len__(self):
        self.len = len(self.data) - self.history_len
        return self.len

    def __getitem__(self, index):
        x_cols = ['Open-N', 'High-N', 'Low-N', 'Close-N', 'Volume-N']
        y_cols = ['Close-N']
        x = self.data.iloc[index: index+self.history_len, :][x_cols].values
        y = self.data.iloc[index+self.history_len, :][y_cols].values.
↪astype('float')
        x = torch.tensor(x).float()
        y = torch.tensor(y).float()
        return x, y

```

```

[21]: # verify dataset instances
train_dataset = MarketDataset(train_data, history_len=48)
print(len(train_dataset))
x, y = train_dataset[21118]
print(x.shape, y.shape)

```

```

21120
torch.Size([48, 5]) torch.Size([1])

```

```

[22]: # test dataset instances
# print(x, y)
# train_data.tail()

```

```

[23]: # test data_loader
train_data_loader = DataLoader(train_dataset, batch_size=4)
X, Y = next(iter(train_data_loader))
print(X.shape, Y.shape)

```

```

torch.Size([4, 48, 5]) torch.Size([4, 1])

```

0.5 5. Define model

```
[24]: class LSTNet(nn.Module):

    def __init__(self):
        super(LSTNet, self).__init__()
        self.num_features = 5
        self.conv1_out_channels = 32
        self.conv1_kernel_height = 7
        self.recc1_out_channels = 64
        self.skip_steps = [4, 24]
        self.skip_reccs_out_channels = [4, 4]
        self.output_out_features = 1
        self.ar_window_size = 7
        self.dropout = nn.Dropout(p = 0.2)

        self.conv1 = nn.Conv2d(1, self.conv1_out_channels,
                                kernel_size=(self.conv1_kernel_height, self.
→num_features))
        self.recc1 = nn.GRU(self.conv1_out_channels, self.recc1_out_channels,
→batch_first=True)
        self.skip_reccs = {}
        for i in range(len(self.skip_steps)):
            self.skip_reccs[i] = nn.GRU(self.conv1_out_channels, self.
→skip_reccs_out_channels[i], batch_first=True)
            self.output_in_features = self.recc1_out_channels + np.dot(self.
→skip_steps, self.skip_reccs_out_channels)
            self.output = nn.Linear(self.output_in_features, self.
→output_out_features)
            if self.ar_window_size > 0:
                self.ar = nn.Linear(self.ar_window_size, 1)

    def forward(self, X):
        """
        Parameters:
        X (tensor) [batch_size, time_steps, num_features]
        """
        batch_size = X.size(0)

        # Convolutional Layer
        C = X.unsqueeze(1) # [batch_size, num_channels=1, time_steps,
→num_features]
        C = F.relu(self.conv1(C)) # [batch_size, conv1_out_channels,
→shrunked_time_steps, 1]
        C = self.dropout(C)
```



```

        C = torch.squeeze(C, 3) # [batch_size, conv1_out_channels,
→shranked_time_steps]

        # Recurrent Layer
        R = C.permute(0, 2, 1) # [batch_size, shrinked_time_steps,
→conv1_out_channels]
        out, hidden = self.recc1(R) # [batch_size, shrinked_time_steps,
→recc_out_channels]
        R = out[:, -1, :] # [batch_size, recc_out_channels]
        R = self.dropout(R)
        #print(R.shape)

        # Skip Recurrent Layers
        shrinked_time_steps = C.size(2)
        for i in range(len(self.skip_steps)):
            skip_step = self.skip_steps[i]
            skip_sequence_len = shrinked_time_steps // skip_step
            # shrinked_time_steps shrinked further
            S = C[:, :, -skip_sequence_len*skip_step:] # [batch_size,
→conv1_out_channels, shrinked_time_steps]
            S = S.view(S.size(0), S.size(1), skip_sequence_len, skip_step) #
→[batch_size, conv1_out_channels, skip_sequence_len,
→skip_step=num_skip_components]
            # note that num_skip_components = skip_step
            S = S.permute(0, 3, 2, 1).contiguous() # [batch_size,
→skip_step=num_skip_components, skip_sequence_len, conv1_out_channels]
            S = S.view(S.size(0)*S.size(1), S.size(2), S.size(3)) #
→[batch_size*num_skip_components, skip_sequence_len, conv1_out_channels]
            out, hidden = self.skip_reccs[i](S) #
→[batch_size*num_skip_components, skip_sequence_len,
→skip_reccs_out_channels[i]]
            S = out[:, -1, :] # [batch_size*num_skip_components,
→skip_reccs_out_channels[i]]
            S = S.view(batch_size, skip_step*S.size(1)) # [batch_size,
→num_skip_components*skip_reccs_out_channels[i]]
            S = self.dropout(S)
            R = torch.cat((R, S), 1) # [batch_size, recc_out_channels +
→skip_reccs_out_channels * num_skip_components]
            #print(S.shape)
            #print(R.shape)

        # Output Layer
        O = F.relu(self.output(R)) # [batch_size, output_out_features=1]

        if self.ar_window_size > 0:
            # set dim3 based on output_out_features

```

```

        AR = X[:, -self.ar_window_size:, 3:4] # [batch_size,
↪ar_window_size, output_out_features=1]
        AR = AR.permute(0, 2, 1).contiguous() # [batch_size,
↪output_out_features, ar_window_size]
        AR = self.ar(AR) # [batch_size, output_out_features, 1]
        AR = AR.squeeze(2) # [batch_size, output_out_features]
        O = O + AR

    return O

```

```

[25]: # test model
model = LSTNet()

for X, Y in train_data_loader:
    print(X.shape)
    out = model(X)
    print(Y.shape, out.shape)
    break

```

```

torch.Size([4, 48, 5])
torch.Size([4, 1]) torch.Size([4, 1])

```

0.6 6. Train model

```

[26]: history_len = 48
batch_size = 8

epochs = 10

lr = 0.01
weight_decay = 0.01

```

```

[27]: train_dataset = MarketDataset(train_data, history_len=history_len)
train_data_loader = DataLoader(train_dataset, batch_size=batch_size,
↪shuffle=True)

model = LSTNet()

criterion = nn.MSELoss()
optimizer = optim.Adam(model.parameters(), lr=lr, weight_decay=weight_decay)

```

```

[27]: train_loss_list = []
for epoch in tqdm(range(epochs)):

    epoch_loss_train = 0
    for i, batch in tqdm(enumerate(train_data_loader, start=1),

```

```

        leave=False, desc="Train",
        total=len(train_data_loader)):

    X, Y = batch
    optimizer.zero_grad()
    Y_pred = model(X)
    loss = criterion(Y_pred, Y)
    loss.backward()
    optimizer.step()

    with open('Log/Running-Loss.txt', 'a+') as file:
        file.write(f'{loss.item()}\n')
    epoch_loss_train += loss.item()

epoch_loss_train = epoch_loss_train / len(train_data_loader)
train_loss_list.append(epoch_loss_train)

with open('Log/Epoch-Loss.txt', 'a+') as file:
    file.write(f'{epoch_loss_train}\n')

```

```
HBox(children=(FloatProgress(value=0.0, max=10.0), HTML(value='')))
```

```
HBox(children=(FloatProgress(value=0.0, description='Train', max=2640.0, style=ProgressStyle(d
```

```
HBox(children=(FloatProgress(value=0.0, description='Train', max=2640.0, style=ProgressStyle(d
```

```
HBox(children=(FloatProgress(value=0.0, description='Train', max=2640.0, style=ProgressStyle(d
```

```
HBox(children=(FloatProgress(value=0.0, description='Train', max=2640.0, style=ProgressStyle(d
```

```
HBox(children=(FloatProgress(value=0.0, description='Train', max=2640.0, style=ProgressStyle(d
```

```
HBox(children=(FloatProgress(value=0.0, description='Train', max=2640.0, style=ProgressStyle(d
```


```
HBox(children=(FloatProgress(value=0.0, description='Train', max=2640.0, style=ProgressStyle(d
```

```
HBox(children=(FloatProgress(value=0.0, description='Train', max=2640.0, style=ProgressStyle(d
```

```
HBox(children=(FloatProgress(value=0.0, description='Train', max=2640.0, style=ProgressStyle(d
```


```
HBox(children=(FloatProgress(value=0.0, description='Train', max=2640.0, style=ProgressStyle(d
```

```
[44]: plt.plot(range(1, len(train_loss_list)+1), train_loss_list, label='Train Loss')
plt.xlabel("Epochs")
plt.legend()
plt.show()
```



```
[29]: with open('Log/Running-Loss.txt', 'r') as file:
      running_loss = pd.Series(file.readlines())
      running_loss_ma = running_loss.rolling(window=100).mean()
      print(len(running_loss_ma))
      plt.plot(running_loss_ma, label='Running Loss')
      plt.xlabel("Iterations")
      plt.show()
```

26400



```
[30]: model_name = f"LSTNet-E{epochs}"
      model_path = f"Models/{model_name}.pth"
      torch.save(model.state_dict(), model_path)
```

0.7 7. Load model

```
[28]: model_name = f"LSTNet-E{epochs}"
      model_path = f"Models/{model_name}.pth"
      model = LSTNet()
      model.load_state_dict(torch.load(model_path))
      model.eval()
```

```
[28]: LSTNet(
  (dropout): Dropout(p=0.2, inplace=False)
  (conv1): Conv2d(1, 32, kernel_size=(7, 5), stride=(1, 1))
  (recc1): GRU(32, 64, batch_first=True)
  (output): Linear(in_features=176, out_features=1, bias=True)
  (ar): Linear(in_features=7, out_features=1, bias=True)
)
```

0.8 8. Make predictions

Predict hourly OHLCV from Jun 1 2020 to Jun 11 2020

```
[29]: data.loc[:, 'Open-N'] = (data['Open'] - min_price) / (max_price- min_price)
data.loc[:, 'High-N'] = (data['High'] - min_price) / (max_price- min_price)
data.loc[:, 'Low-N'] = (data['Low'] - min_price) / (max_price- min_price)
data.loc[:, 'Close-N'] = (data['Close'] - min_price) / (max_price- min_price)
data.loc[:, 'Volume-N'] = (data['Volume'] - min_volume) / (max_volume-  

↪min_volume)
data.head()
```

```
[29]:      Open-Time      Open      High      Low      Close      Volume \
0 2018-01-01 00:00:00  0.000623  0.000625  0.000614  0.000618  130791.0
1 2018-01-01 01:00:00  0.000618  0.000625  0.000611  0.000622   89072.0
2 2018-01-01 02:00:00  0.000622  0.000625  0.000620  0.000622  101296.0
3 2018-01-01 03:00:00  0.000622  0.000624  0.000619  0.000622   83701.0
4 2018-01-01 04:00:00  0.000622  0.000628  0.000622  0.000625  103000.0

      Open-N      High-N      Low-N      Close-N      Volume-N
0  0.021938  0.022318  0.019779  0.020873  0.025258
1  0.020873  0.022388  0.019193  0.021767  0.017201
2  0.021770  0.022493  0.021181  0.021681  0.019562
3  0.021679  0.022148  0.020932  0.021809  0.016164
4  0.021753  0.023107  0.021676  0.022391  0.019891
```

```
[30]: batch_size = 16
dataset = MarketDataset(data, history_len=history_len)
data_loader = DataLoader(dataset, batch_size=batch_size, shuffle=False)
```

```
[31]: for i, batch in tqdm(enumerate(data_loader, start=1), leave=False,  

↪total=len(data_loader)):

    X, Y = batch
    Y_pred = model(X).detach().numpy()
    if i == 1:
        predictions = Y_pred
    else:
        predictions = np.concatenate((predictions, Y_pred), axis=0)

    #predictions[:, 0:4] = predictions[:, 0:4] * (max_price - min_price) + min_price
    #predictions[:, 4] = predictions[:, 4] * (max_volume - min_volume) + min_volume
    predictions = predictions * (max_price - min_price) + min_price

    #columns = ['Open', 'High', 'Low', 'Close', 'Volume']
    columns = ['Close']
    predictions = pd.DataFrame(predictions, columns=columns)
    print(predictions.shape)
    predictions.head()
```

```
HBox(children=(FloatProgress(value=0.0, max=1336.0), HTML(value='')))
```

```
(21370, 1)
```

```
[31]:      Close
0  0.000621
1  0.000616
2  0.000617
3  0.000612
4  0.000610
```

```
[32]: print(data.shape)
      print(predictions.shape)
```

```
(21418, 11)
(21370, 1)
```

```
[33]: cols = ['Open-Time', 'Open', 'High', 'Low', 'Close', 'Volume']
      data = data[cols]
      data.tail()
```

```
[33]:      Open-Time      Open      High      Low      Close      Volume
21413 2020-06-11 05:00:00  0.001763  0.001767  0.001762  0.001762  29215.63
21414 2020-06-11 06:00:00  0.001762  0.001769  0.001760  0.001766  18447.87
21415 2020-06-11 07:00:00  0.001766  0.001769  0.001764  0.001769  49849.23
21416 2020-06-11 08:00:00  0.001769  0.001771  0.001767  0.001770  39773.38
21417 2020-06-11 09:00:00  0.001770  0.001771  0.001768  0.001770   3454.59
```

```
[34]: predictions['Open-Time'] = pd.Series(predictions.index).apply(lambda x:
    ↪data['Open-Time'][x+history_len])
      cols = ['Open-Time', 'Close']
      predictions = predictions[cols]
      predictions.tail()
```

```
[34]:      Open-Time      Close
21365 2020-06-11 05:00:00  0.001759
21366 2020-06-11 06:00:00  0.001759
21367 2020-06-11 07:00:00  0.001760
21368 2020-06-11 08:00:00  0.001761
21369 2020-06-11 09:00:00  0.001762
```

0.9 9. Visualise Predictions

```
[35]: min_time = datetime.strptime('2020-06-01 00:00:00', '%Y-%m-%d %H:%M:%S')
      max_time = datetime.strptime('2020-06-10 23:00:00', '%Y-%m-%d %H:%M:%S')
      data_mask = (data['Open-Time'] >= min_time) & (data['Open-Time'] <= max_time)
      predictions_mask = (predictions['Open-Time'] >= min_time) &
    ↪(predictions['Open-Time'] <= max_time)
```

```
print(data_mask.sum(), predictions_mask.sum())
```

240 240

```
[36]: actuals_df = data[data_mask].copy().set_index('Open-Time', drop=True)
      predictions_df = predictions[predictions_mask].copy().set_index('Open-Time',
      ↪drop=True)
```

```
[37]: actuals_df.head()
```

```
[37]:
```

	Open	High	Low	Close	Volume
Open-Time					
2020-06-01 00:00:00	0.001809	0.001813	0.001800	0.001812	73009.76
2020-06-01 01:00:00	0.001811	0.001826	0.001808	0.001824	54387.68
2020-06-01 02:00:00	0.001824	0.001834	0.001822	0.001834	64901.54
2020-06-01 03:00:00	0.001834	0.001836	0.001822	0.001826	55693.64
2020-06-01 04:00:00	0.001825	0.001825	0.001814	0.001815	50167.58

```
[38]: predictions_df.head()
```

```
[38]:
```

	Close
Open-Time	
2020-06-01 00:00:00	0.001818
2020-06-01 01:00:00	0.001815
2020-06-01 02:00:00	0.001815
2020-06-01 03:00:00	0.001816
2020-06-01 04:00:00	0.001815

```
[39]: lines = []

      line = {
          'x':actuals_df.index,
          'y':actuals_df['Close'],
          'name':'Close-Actuals'
      }
      lines.append(line)

      line = {
          'x':predictions_df.index,
          'y':predictions_df['Close'],
          'name':'Close-Predictions'
      }
      lines.append(line)
```

```
[40]: fig = go.Figure(lines)
      iplot(fig)
```



```
[42]: from IPython.display import Image
      Image(filename='Prediction-Visualisation.png')
```

[42]:



[]: